

NxM Templates for SuperCDMS SNOLAB Run 4 — Ge Activation Data

K-line event selection · free-pretrigger two-exponential fits (*future step: three-, four-exponential fits*) · 1x1
and NxM (PCA) templates


Table of Contents

1. **Dataset:** Ge activation K-line event selection, raw trace cache (13 detectors)
2. **Pulse fit:** 100 kHz low-pass, free-pretrigger two-exponential fit, alignment at 16050
3. **Data cleaning:** fit_ok, NRMSE valley cut (0.4), rise-time ceiling (0.3 ms)
4. **Cross-checks:** cuts remove noise, real pulses stay; slow-fall follow-up
5. **Template generation:** 1x1 weighted template, NxM PCA templates
6. **Template file and future steps:** cdmsbats ROOT files, ready for merge; three-, four-exponential fits
7. **Backup:** weak detectors, rise-time cut details, full results for every zip

Ge activation dataset - R4

Credit: Prof. Tarek Saab

- Per detector, Prof. Saab's study marked the 10.37 keV K-line position in **PTOFamps** (the red line in each panel below)
- A **PTOFamps** window bracketing the red line: position $\times/\div 1.35$, a rough eyeballed choice; deliberately minimal, no other cuts
- For every selected event: cache the fully unprocessed raw MIDAS traces of all channels + event metadata, saved as per-series pickle files (reusable for any later analysis)
- Events in the window per zip: Z1 10.1k Z4 26.1k Z6 17.0k Z7 2.2k Z9 2.6k Z10 7.6k Z13 9.1k Z15 4.2k Z16 1.6k Z18 22.1k Z19 23.3k Z22 25.0k Z24 22.9k (total 174k)
- **This talk focuses on Z7, the best detector; all other detectors are covered in the backup slides**




All 13 detectors, PTOFamps spectra (Prof. Saab's study)


Fit quality at a glance: event × channel grids

- Real events and noise triggers can be told apart by visual comparison (no cut applied at this stage)

NOISE



REAL PULSE



Per-trace algorithm: low-pass → fit → align

1. **Low-pass** — 100 kHz Butterworth
2. **Normalize** — subtract the baseline (median of the pre-pulse samples), scale the trace to unit peak
3. **Fit** — two-exponential model:

$$y(t) = A[e^{-(t - t_0)/\tau_{\text{fall}}} - e^{-(t - t_0)/\tau_{\text{rise}}}] + b$$


t_0 = pretrigger, free within 16050 ± 3000 samples (τ_{rise} , τ_{fall} : rise / fall time)

Fit parameters: A amplitude · τ_{rise} rise time constant · τ_{fall} fall time constant · t_0 pulse onset (pretrigger) · b baseline

4. **Two quality checks per trace** (calculated and recorded here, not cut yet):

- **fit_ok** — amplitude is positive, and the pulse rises faster than it falls
- **NRMSE** — RMS of the fit residual ÷ pulse peak (smaller = better fit)


5. Align



Z7 PBS1: every fitted event drawn at the common pretrigger, no cuts of any kind (n = 2008 events)


Quality cut 1: $\text{NRMSE} \leq 0.4$ — where the number comes from

- NRMSE of fit_ok events is bimodal: good fits (median ≈ 0.05 – 0.1) vs noise triggers (≈ 1 – 2), valley at ≈ 0.4 – 0.5
- The cut sits in the valley: eyeball read off the distribution
- Worse detectors (Z1, Z4, Z6, Z18, Z19, Z22, Z24): same bimodal picture, noise peak dominates




Z7 PBS1: two clean populations, log-log axes

Aligned fitted curves: before and after the NRMSE cut



All fit_ok fitted curves (Z7 PBS1), 200-curve sample drawn




After $\text{NRMSE} \leq 0.4$ (Z7 PBS1): one tight shape family

The rejected population is noise — verified in the raw traces

- Example of NRMSE-rejected events (median NRMSE > 0.4 across channels): the raw traces show no pulse
- the cut removes noise triggers, not physics


Rejected events: raw vs fit



Z7: NRMSE-rejected events — raw traces are noise

Aligned curves: kept vs rejected


median NRMSE: **kept 0.045 (n = 1931)** · **rejected 1.873 (n = 77)**



Z7 PBS1 — **passes the cut (kept)** — **cut away = rejected**
faint blue = measured traces 1931 kept / 77 cut (~3.8%)


Slow-fall tail in one particular channel

- The post-cut plot still shows slow-fall tails → sample them: $\text{NRMSE} \leq 0.4$ AND $\tau_{\text{fall}} > 1.5$ ms, 10 random events, raw vs fit drawn in all 12 channels
- The sampled events are real pulses. Their extreme fall times trace to one channel: only PDS2 swings (τ_{fall} median 0.51 ms vs ≈ 0.25 ms elsewhere)
- Conclusion: no τ_{fall} cut — slow-fall events passing NRMSE stay in
- PDS2 alone: can not make useful templates; hard to fix in analysis/software, the noise itself needs to be fixed
- Temporary solution: use the PDS1 template in place of PDS2 (applied in the delivered ROOT files)




Z7, 3 sampled slow-fall events: normal in PBS2/PCS2/PES2, swings only in PDS2

Fitted τ_{fall} distribution, same 0–7 ms axis:




Z7 PAS1 — normal channel: narrow
(median 0.25 ms)



Z7 PDS2 — bad channel: broad tail
(median 0.51 ms)

Results: 1x1 Templates


- NRMSE-weighted mean of the fit_ok 2-exp curves
- Smooth & noise-free by construction




Z7 PBS1 — blue: the fitted curves (200 of 1931 drawn); red: the 1x1 template

Results: NxM Templates


- channel-specific PCA over the fitted curves ($\text{fit_ok} + \text{NRMSE} \leq 0.4$)
- nxm0 = mean shape; nxm1 – 4 = principal components; $\text{real pulse} = \sum_i \text{amp}_i \cdot \text{nxm}_i$
- $\text{PC1} + \text{PC2} \approx 96\text{--}98\%$ of the shape variance



NRMSE-weighted mean of all fit_ok curves ($n = 2008$) — exactly the 1x1 template



Plain mean of the PCA input curves ($n = 1931$) — exactly nxm0



Z7 PBS1: nxm0 (mean) + nxm1 – 4 (PCs) — the oscillating components encode rise/fall-time variation

Template file


- Two template sets for all 13 detectors: 1x1 (2-exp weighted) + NxM PCA
- In cdmsbats_config feature branch, ready for merge

Future Steps


- Three-, four-exponential fits
- Exercise NxM processing
- Other detectors: worse noise, a few more cuts added, in the backup slides

Backup — weak detectors (example: Z22)


- Z1 / Z4 / Z6 / Z18 / Z19 / Z22 / Z24: K-line sits inside the noise population — the PTOF window admits a noise-dominated mixture
- Same bimodal NRMSE picture, noise peak dominates; the same 0.4 cut digs the real pulses out — Z22 PCS1: 7198 kept / 4788 cut (~40%)
- Kept bundle broader than on quiet detectors; steep red curves = fits latching onto sharp noise spikes, not fast pulses being cut
- $\tau_{\text{rise}} \leq 0.3$ ms (after NRMSE): removes 55–74% on weak zips (Z22: 71%) vs 1.6% on Z7 — residual slow drift



Z22 PAS1: NRMSE histogram — noise dominates




Z22 PCS1: kept (green) vs cut (red)



Z22 PCS1 after $\text{NRMSE} \leq 0.4 + \tau_{\text{rise}} \leq 0.3$ ms — final shape family (760 fits)

Backup — the $\tau_{\text{rise}} \leq 0.3$ ms ceiling (PCA input)


- Slow baseline drift survives NRMSE (a slow 2-exp hugs it) → mimics a slow rise
- Fast pulses at $\tau_{\text{rise}} \approx 0.1$ ms (p90 ≈ 0.15 ms) → ceiling at 0.3 ms blocks the drift tail
- Removed after the NRMSE cut: Z7 1.6% (quiet 2–6%) — weak 55–74% (Z22 71%) — all zips 54%
- Trade-off: trims the very slowest genuine pulses → decision pending




Z7 PAS1: fitted τ_{rise} (median 0.105 ms) and τ_{fall} distributions

Backup — τ_{rise} cut: a good vs a bad channel (Z7)


- Good channel PBS1: tight clean bundle — the $\tau_{\text{rise}} \leq 0.3$ ms cut removes $\sim 0\%$ (nearly free on quiet channels)
- Bad channel PDS2: broad, messy fan — the channel itself is bad; the cut trims the slow-rise curves, 1462 $\rightarrow$ 1154 (21%)




PBS1 (good) — $\text{NRMSE} \leq 0.4$



PBS1 (good) — $+\tau_{\text{rise}} \leq 0.3$ ms (unchanged)




PDS2 (bad) — $\text{NRMSE} \leq 0.4$



PDS2 (bad) — $+\tau_{\text{rise}} \leq 0.3$ ms (slow risers trimmed)

Backup — what the τ_{rise} cut removes: noise NRMSE missed (Z22)

- Exactly the events that PASS $\text{NRMSE} \leq 0.4$ but are REMOVED by $\tau_{\text{rise}} \leq 0.3$ ms — the noise 0.4 let through, 0.3 catches
- PCS1 / PDS1 pass at $\text{NRMSE} \approx 0.15$ (a slow 2-exp hugs a slow baseline drift) — but the raw trace has no clean fast pulse
- Trade-off: the same cut also trims genuine slow-rise pulses — documented, decision pending



Z22: 3 events $\times$ PAS1/PBS1/PCS1/PDS1 — pass $\text{NRMSE} \leq 0.4$, removed by $\tau_{\text{rise}} \leq 0.3$ ms; PCS1/PDS1 are pure slow drift

Backup — Z7 full results gallery


every diagnostic and result figure of the Z7 pipeline, in processing order

- Dataset: the K-line selection window
- Example grids: LP + fit and raw vs fit, event × channel
- Fit diagnostics: NRMSE, fitted pretrigger, τ_{rise} / τ_{fall} distributions
- Aligned fans: no cut $\rightarrow$ $\text{NRMSE} \leq 0.4 \rightarrow + \tau_{\text{rise}} \leq 0.3$ ms; kept vs rejected; measured traces
- Special populations: rejected events, well-fit slow-rise events, PDS2 slow-fall study
- Templates: PCA nxm0–4 per channel, summed PT / PS1 / PS2
- Tall figures are cut at white space into full-width panels, two per page; grids into 3-event × half-channel blocks

Backup — Z7 · K-line selection window

figure: kline_zip7.png


Summed PTOFamps spectrum of Z7; red line = Prof. Saab's 10.37 keV K-line position, shaded band = the $\times/\div 1.35$ selection window; everything inside is cached raw.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (1/10)

figure: zip7_fit_examples.png — event rows 1–3 of 15, left half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (2/10)

figure: zip7_fit_examples.png — event rows 1–3 of 15, right half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (3/10)

figure: zip7_fit_examples.png — event rows 4–6 of 15, left half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (4/10)

figure: zip7_fit_examples.png — event rows 4–6 of 15, right half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (5/10)

figure: zip7_fit_examples.png — event rows 7–9 of 15, left half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (6/10)

figure: zip7_fit_examples.png — event rows 7–9 of 15, right half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (7/10)

figure: zip7_fit_examples.png — event rows 10–12 of 15, left half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (8/10)

figure: zip7_fit_examples.png — event rows 10–12 of 15, right half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (9/10)

figure: zip7_fit_examples.png — event rows 13–15 of 15, left half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · fit_examples — LP trace vs 2-exp fit (10/10)

figure: zip7_fit_examples.png — event rows 13–15 of 15, right half of the channels


First 15 cached events × all channels; blue = 100 kHz LP trace, red = 2-exp fit. A real event fits well in every channel at once; a noise trigger fails across the whole row.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (1/10)

figure: zip7_raw_vs_fit_examples.png — event rows 1–3 of 15, left half of the channels


Same 15 events × all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (2/10)

figure: zip7_raw_vs_fit_examples.png — event rows 1–3 of 15, right half of the channels


Same 15 events × all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (3/10)

figure: zip7_raw_vs_fit_examples.png — event rows 4–6 of 15, left half of the channels


Same 15 events $\times$ all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (4/10)

figure: zip7_raw_vs_fit_examples.png — event rows 4–6 of 15, right half of the channels


Same 15 events $\times$ all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (5/10)

figure: zip7_raw_vs_fit_examples.png — event rows 7–9 of 15, left half of the channels


Same 15 events × all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (6/10)

figure: zip7_raw_vs_fit_examples.png — event rows 7–9 of 15, right half of the channels


Same 15 events $\times$ all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (7/10)

figure: zip7_raw_vs_fit_examples.png — event rows 10–12 of 15, left half of the channels


Same 15 events × all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (8/10)

figure: zip7_raw_vs_fit_examples.png — event rows 10–12 of 15, right half of the channels


Same 15 events $\times$ all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (9/10)

figure: zip7_raw_vs_fit_examples.png — event rows 13–15 of 15, left half of the channels


Same 15 events × all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · raw_vs_fit_examples — raw trace vs fit (10/10)

figure: zip7_raw_vs_fit_examples.png — event rows 13–15 of 15, right half of the channels


Same 15 events × all channels; gray = raw unfiltered trace, blue = LP, red = fit, per-panel NRMSE top-right. Noise vs pulse size is directly visible in the raw data.



Backup — Z7 · nrmse — NRMSE distribution per channel (1/6)

figure: zip7_nrmse.png — channel panels 1–2 of 11, top to bottom


NRMSE = RMS(fit residual)/fitted peak, log-log. Bimodal: good fits ($\sim 0.05\text{--}0.1$) vs noise triggers ($\sim 1\text{--}2$); the valley at ≈ 0.4 sets quality cut 1.



Backup — Z7 · nrmse — NRMSE distribution per channel (2/6)

figure: zip7_nrmse.png — channel panels 3–4 of 11, top to bottom


NRMSE = RMS(fit residual)/fitted peak, log-log. Bimodal: good fits ($\sim 0.05\text{--}0.1$) vs noise triggers ($\sim 1\text{--}2$); the valley at ≈ 0.4 sets quality cut 1.



Backup — Z7 · nrmse — NRMSE distribution per channel (3/6)

figure: zip7_nrmse.png — channel panels 5–6 of 11, top to bottom


NRMSE = RMS(fit residual)/fitted peak, log-log. Bimodal: good fits ($\sim 0.05\text{--}0.1$) vs noise triggers ($\sim 1\text{--}2$); the valley at ≈ 0.4 sets quality cut 1.



Backup — Z7 · nrmse — NRMSE distribution per channel (4/6)

figure: zip7_nrmse.png — channel panels 7–8 of 11, top to bottom


NRMSE = RMS(fit residual)/fitted peak, log-log. Bimodal: good fits ($\sim 0.05\text{--}0.1$) vs noise triggers ($\sim 1\text{--}2$); the valley at ≈ 0.4 sets quality cut 1.



Backup — Z7 · nrmse — NRMSE distribution per channel (5/6)

figure: zip7_nrmse.png — channel panels 9–10 of 11, top to bottom


NRMSE = RMS(fit residual)/fitted peak, log-log. Bimodal: good fits ($\sim 0.05\text{--}0.1$) vs noise triggers ($\sim 1\text{--}2$); the valley at ≈ 0.4 sets quality cut 1.



Backup — Z7 · nrmse — NRMSE distribution per channel (6/6)

figure: zip7_nrmse.png — channel panels 11–11 of 11, top to bottom

NRMSE = RMS(fit residual)/fitted peak, log-log. Bimodal: good fits ($\sim 0.05\text{--}0.1$) vs noise triggers ($\sim 1\text{--}2$); the valley at ≈ 0.4 sets quality cut 1.




Backup — Z7 · pretrigger — fitted onset distribution per channel

(1/6)

The onset is a FRE fit parameter; fitted values cluster ~230 samples after the nominal 16050 — pinning it inside the fit would bias every other parameter.

Figure: z1p7_pretrigger.png — channel pairs 1-2 at 11, top to bottom




Backup — Z7 · pretrigger — fitted onset distribution per channel

(2/6)

The onset is a FREE fit parameter; fitted values cluster ~230 samples after the nominal 16050 — pinning it inside the fit would bias every other parameter.

Figure: Z7_pretrigger.png — channel pairs 3-4 at 11, top to bottom




Backup — Z7 · pretrigger — fitted onset distribution per channel

(3/6)

The onset is a FREE fit parameter; fitted values cluster ~230 samples after the nominal 16050 — pinning it inside the fit would bias every other parameter.

Figure: Z7_pretrigger.png — channel pairs 5-6 at 11, top to bottom





Backup — Z7 · pretrigger — fitted onset distribution per channel

(4/6)

The onset is a FREE fit parameter; fitted values cluster ~230 samples after the nominal 16050 — pinning it inside the fit would bias every other parameter.

Figure: Z7_pretrigger.png — channel pairs 7-9 at 11, top to bottom




Backup — Z7 · pretrigger — fitted onset distribution per channel

(5/6)

The onset is a FREE fit parameter: fitted values cluster ~230 samples after the nominal 16050 — pinning it inside the fit would bias every other parameter.

Figure: Z7_pretrigger.png — channel pairs 9-10 only, top to bottom




Backup — Z7 · pretrigger — fitted onset distribution per channel

(6/6)

The onset is a FRE fit parameter: fitted values cluster ~230 samples after the nominal 16050 — pinning it inside the fit would bias every other parameter.

Figure: Z7_pretrigger.png — channel pairs 11-11 of 11, top to bottom




Backup — Z7 · time_constants — fitted τ_{rise} / τ_{fall} per channel

(1/6)

Fast pulse population at $\tau_{\text{rise}} = 0.1$ ms; isolated spikes at the parameter bounds are noise fits pinned at the fit limits, removed by the NRMSE cut.

Figure: Z7_time_constants.png — channel panels 1–2 of 11, top to bottom




Backup — Z7 · time_constants — fitted τ_{rise} / τ_{fall} per channel

(2/6)

Fast-rising population at $\tau_{\text{rise}} = 0.1$ ms; isolated spikes at the parameter bounds are noise fits pinned at the fit limits, removed by the NRMSE cut.

Figure: Z7_time_constants.png — channel panels B–4 or 11, top to bottom




Backup — Z7 · time_constants — fitted τ_{rise} / τ_{fall} per channel

(3/6)

Fast-rising population at $\tau_{\text{rise}} = 0.1$ ms; isolated spikes at the parameter bounds are noise fits pinned at the fit limits, removed by the NRMSE cut.


Figure: Z7_time_constants.png — channel panels D–G of 11, top to bottom



Backup — Z7 · time_constants — fitted τ_{rise} / τ_{fall} per channel

(4/6)


Fast pulse population at $\tau_{\text{rise}} = 0.1$ ms; isolated spikes at the parameter bounds are noise fits pinned at the fit limits, removed by the NRMSE cut.



Backup — Z7 · time_constants — fitted τ_{rise} / τ_{fall} per channel

(5/6)


Fast pulse population at $\tau_{\text{rise}} = 0.1$ ms; isolated spikes at the parameter bounds are noise fits pinned at the fit limits, removed by the NRMSE cut.



Backup — Z7 · time_constants — fitted τ_{rise} / τ_{fall} per channel (6/6)

Fast pulse population at $\tau_{\text{rise}} = 0.1$ ms; isolated spikes at the parameter bounds are noise fits pinned at the fit limits, removed by the NRMSE cut.


Figure: Z7_time_constants.png — channel panels 11–11 of 11, top to bottom



Backup — Z7 · fitted_curves_overlay — fan, no cut (1/6)

figure: zip7_fitted_curves_overlay.png — channel panels 1–2 of 11, top to bottom


Every fit_ok fitted curve at the common pretrigger 16050, peak-normalized. Pure shape distribution; the broad slow curves are noise fits that the cuts remove.



Backup — Z7 · fitted_curves_overlay — fan, no cut (2/6)

figure: zip7_fitted_curves_overlay.png — channel panels 3–4 of 11, top to bottom


Every fit_ok fitted curve at the common pretrigger 16050, peak-normalized. Pure shape distribution; the broad slow curves are noise fits that the cuts remove.



Backup — Z7 · fitted_curves_overlay — fan, no cut (3/6)

figure: zip7_fitted_curves_overlay.png — channel panels 5–6 of 11, top to bottom


Every fit_ok fitted curve at the common pretrigger 16050, peak-normalized. Pure shape distribution; the broad slow curves are noise fits that the cuts remove.



Backup — Z7 · fitted_curves_overlay — fan, no cut (4/6)

figure: zip7_fitted_curves_overlay.png — channel panels 7–8 of 11, top to bottom


Every fit_ok fitted curve at the common pretrigger 16050, peak-normalized. Pure shape distribution; the broad slow curves are noise fits that the cuts remove.



Backup — Z7 · fitted_curves_overlay — fan, no cut (5/6)

figure: zip7_fitted_curves_overlay.png — channel panels 9–10 of 11, top to bottom


Every fit_ok fitted curve at the common pretrigger 16050, peak-normalized. Pure shape distribution; the broad slow curves are noise fits that the cuts remove.



Backup — Z7 · fitted_curves_overlay — fan, no cut (6/6)

figure: zip7_fitted_curves_overlay.png — channel panels 11–11 of 11, top to bottom


Every fit_ok fitted curve at the common pretrigger 16050, peak-normalized. Pure shape distribution; the broad slow curves are noise fits that the cuts remove.



Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4$ (1/6)

figure: zip7_fitted_curves_overlay_nrmse0.4.png — channel panels 1–2 of 11, top to bottom


Same fan after quality cut 1: the noise fan collapses, the fast-pulse bundle remains.



Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4$ (2/6)

figure: zip7_fitted_curves_overlay_nrmse0.4.png — channel panels 3–4 of 11, top to bottom


Same fan after quality cut 1: the noise fan collapses, the fast-pulse bundle remains.



Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4$ (3/6)

figure: zip7_fitted_curves_overlay_nrmse0.4.png — channel panels 5–6 of 11, top to bottom


Same fan after quality cut 1: the noise fan collapses, the fast-pulse bundle remains.



Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4$ (4/6)

figure: zip7_fitted_curves_overlay_nrmse0.4.png — channel panels 7–8 of 11, top to bottom


Same fan after quality cut 1: the noise fan collapses, the fast-pulse bundle remains.



Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4$ (5/6)

figure: zip7_fitted_curves_overlay_nrmse0.4.png — channel panels 9–10 of 11, top to bottom


Same fan after quality cut 1: the noise fan collapses, the fast-pulse bundle remains.



Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4$ (6/6)


figure: zip7_fitted_curves_overlay_nrmse0.4.png — channel panels 11–11 of 11, top to bottom

Same fan after quality cut 1: the noise fan collapses, the fast-pulse bundle remains.




Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4 + \tau_{\text{rise}} \leq 0.3 \text{ ms}$ (1/6)

The exact PCA input population: the clean fast-pulse bundle only
Figure: zip_fitted_curves_overlay_insec.4_insec.30ms.png — channel panels 1–2 of 11, top to bottom




Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4 + \tau_{\text{rise}} \leq 0.3 \text{ ms}$ (2/6)

The exact PCA input population: the clean fast-pulse bundle only
Figure: zip7_fitted_curves_overlay_insec.4_insec.30ms.png — channel panels 3–4 of 11, top to bottom




Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4 + \tau_{\text{rise}} \leq 0.3 \text{ ms}$ (3/6)

The exact PCA input population: the clean fast-pulse bundle only
Figure: z7_fitted_curves_overlay_insec.4_insec.30ms.png — channel panels 5–6 of 11, top to bottom




Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4 + \tau_{\text{rise}} \leq 0.3 \text{ ms}$ (4/6)

The exact PCA input population: the clean fast-pulse bundle only
Figure: z7_fitted_curves_overlay_insec.4_insec.30ms.png — channel panels 7–8 of 11, top to bottom




Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4 + \tau_{\text{rise}} \leq 0.3 \text{ ms}$ (5/6)

The exact PCA input population: the clean fast-pulse bundle only
Figure: zip_fitted_curves_overlay_insec.4_insec.50ms.png — channel panels 9–10 of 11, top to bottom



Backup — Z7 · fitted_curves_overlay — after $\text{NRMSE} \leq 0.4 + \tau_{\text{rise}} \leq 0.3 \text{ ms}$ (6/6)


The exact PCA fitted population: the clean fast-pulse bundle only
Figure: Z7_fitted_curves_overlay_insec.4_insec.30ms.png — channel panels 11–11 of 11, top to bottom



Backup — Z7 · overlay_fan_cut — kept vs rejected fits (1/6)

figure: zip7_overlay_fan_cut_nrmse0.4.png — channel panels 1–2 of 11, top to bottom


Gray-blue = aligned measured traces; green = fits with $\text{NRMSE} \leq 0.4$ (kept); red = $\text{NRMSE} > 0.4$ (cut away). Median NRMSE and counts of both populations stamped per channel.



Backup — Z7 · overlay_fan_cut — kept vs rejected fits (2/6)

figure: zip7_overlay_fan_cut_nrmse0.4.png — channel panels 3–4 of 11, top to bottom


Gray-blue = aligned measured traces; green = fits with $\text{NRMSE} \leq 0.4$ (kept); red = $\text{NRMSE} > 0.4$ (cut away). Median NRMSE and counts of both populations stamped per channel.



Backup — Z7 · overlay_fan_cut — kept vs rejected fits (3/6)

figure: zip7_overlay_fan_cut_nrmse0.4.png — channel panels 5–6 of 11, top to bottom


Gray-blue = aligned measured traces; green = fits with $\text{NRMSE} \leq 0.4$ (kept); red = $\text{NRMSE} > 0.4$ (cut away). Median NRMSE and counts of both populations stamped per channel.



Backup — Z7 · overlay_fan_cut — kept vs rejected fits (4/6)

figure: zip7_overlay_fan_cut_nrmse0.4.png — channel panels 7–8 of 11, top to bottom


Gray-blue = aligned measured traces; green = fits with $\text{NRMSE} \leq 0.4$ (kept); red = $\text{NRMSE} > 0.4$ (cut away). Median NRMSE and counts of both populations stamped per channel.



Backup — Z7 · overlay_fan_cut — kept vs rejected fits (5/6)

figure: zip7_overlay_fan_cut_nrmse0.4.png — channel panels 9–10 of 11, top to bottom


Gray-blue = aligned measured traces; green = fits with $\text{NRMSE} \leq 0.4$ (kept); red = $\text{NRMSE} > 0.4$ (cut away). Median NRMSE and counts of both populations stamped per channel.



Backup — Z7 · overlay_fan_cut — kept vs rejected fits (6/6)

figure: zip7_overlay_fan_cut_nrmse0.4.png — channel panels 11–11 of 11, top to bottom





Gray-blue = aligned measured traces; green = fits with $\text{NRMSE} \leq 0.4$ (kept); red = $\text{NRMSE} > 0.4$ (cut away). Median NRMSE and counts of both populations stamped per channel.



Backup — Z7 · aligned_overlay — shift-aligned measured traces

(1/6)





Blue = measured LP traces shifted by (fitted pre-trigger - 16050); red = mean of all fit_ok events; orange = NRMSE-weighted mean of the fitted curves. Rise edges line up at 16050.



Backup — Z7 · aligned_overlay — shift-aligned measured traces

(2/6)





Blue = measured LP traces shifted by (fitted pre-trigger - 16050); red = mean of all fit_ok events; orange = NRMSE-weighted mean of the fitted curves. Rise edges line up at 16050.



Backup — Z7 · aligned_overlay — shift-aligned measured traces





(3/6)

Blue = measured LP traces shifted by (fitted pre-trigger - 16050); red = mean of all fit_ok events; orange = NRMSE-weighted mean of the fitted curves. Rise edges line up at 16050.



Backup — Z7 · aligned_overlay — shift-aligned measured traces (4/6)

Blue = measured LP traces shifted by (fitted pre-trigger = 16050), red = mean of all fit_ok events, orange = NRMSE-weighted mean of the fitted curves. Rise edges line up at 16050.







Backup — Z7 · aligned_overlay — shift-aligned measured traces

(5/6)



Blue = measured LP traces shifted by (fitted pre-trigger - 16050); red = mean of all fit_ok events; orange = NRMSE-weighted mean of the fitted curves. Rise edges line up at 16050.

Figure: lp_aligned_overlay.png — channel panels 9–10 on 11, top to bottom



Backup — Z7 · aligned_overlay — shift-aligned measured traces (6/6)

Blue = measured LP traces shifted by (fitted pre-trigger - 16050); red = mean of all fit events; orange = NRMSE-weighted mean of the fitted curves. Rise edges line up at 16050.




Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(1/10)

Events with median NRMSE > 0.4 across channels: the raw traces show no pulse in any channel — the rejected population is noise triggers, not physics.

Figure: Z7_slow_rise_events.png — event rows 1-6 of 10, left half of the channels




Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(2/10)

Events with median NRMSE > 0.4 across channels: the raw traces show no pulse in any channel — the rejected population is noise triggers, not physics.


Figure: Z7_slow_rise_events.png — event rows 1-6 of 19, right half of the channels



Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(3/10)

Events with median NRMSE > 0.4 across channels: the raw traces show no pulse in any channel — the rejected population is noise triggers, not physics.




Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(4/10)

Events with median NRMSE > 0.4 across channels: the raw traces show no pulse in any channel — the rejected population is noise triggers, not physics.


Figure: Z7_slow_rise_events.png — event rows 4-6 of 19, right half of the channels



Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(5/10)

Events with median NRMSE > 0.4 across channels, the raw traces show no pulse in any channel — the rejected population is noise triggers, not physics.




Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(6/10)

Events with median NRMSE > 0.4 across channels: the raw traces show no pulse in any channel — the rejected population is noise triggers, not physics.


Figure: Z7_slow_rise_events.png — event rows 1-9 of 19, right half of the channels



Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(7/10)


Events with median NRMSE > 0.4 across channels, the raw traces show no pulse in any channel, the rejected population is noise triggers, not physics.



Backup — Z7 · slow_rise_events — the NRMSE-rejected population (8/10)

Events with median NRMSE > 0.4 across channels: the raw traces show no rise in any channel — the rejected population is noise triggers, not physics.


Figure: Z7_slow_rise_events.png — event rows 10–12 of 15, right half of the channels



Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(9/10)

Events with median NRMSE > 0.4 across channels: the raw traces show no pulse in any channel; the rejected population is noise triggers, not physics.




Backup — Z7 · slow_rise_events — the NRMSE-rejected population

(10/10)

Events with median NRMSE > 0.4 across channels; the raw traces show no pulse in any channel — the rejected population is noise triggers, not physics.


Figure: Z7_slow_rise_events.png — event rows 13–15 of 15, right half of the channels



Backup — Z7 · shadow_events — well-fit slow-rise events (1/2)

figure: zip7_shadow_events.png — event rows 1–3 of 3, left half of the channels


Median NRMSE ≤ 0.4 AND median $\tau_{\text{rise}} > 0.2$ ms: genuinely slow, well-fit pulses — kept, and exactly what the NxM multi-templates are for.



Backup — Z7 · shadow_events — well-fit slow-rise events (2/2)

figure: zip7_shadow_events.png — event rows 1–3 of 3, right half of the channels


Median NRMSE ≤ 0.4 AND median $\tau_{\text{rise}} > 0.2$ ms: genuinely slow, well-fit pulses — kept, and exactly what the NxM multi-templates are for.



Backup — Z7 · slow_fall_events — long- τ _fall study (PDS2) (1/8)

figure: zip7_PDS2_slow_fall.png — event rows 1–3 of 10, left half of the channels


10 random events with $\tau_{\text{fall}} > 1.5$ ms drawn in every channel: 11 channels show normal fast pulses; the long fall is the PDS2 low-frequency hardware noise, not slow physics.



Backup — Z7 · slow_fall_events — long- τ _fall study (PDS2) (2/8)

figure: zip7_PDS2_slow_fall.png — event rows 1–3 of 10, right half of the channels


10 random events with $\tau_{\text{fall}} > 1.5$ ms drawn in every channel: 11 channels show normal fast pulses; the long fall is the PDS2 low-frequency hardware noise, not slow physics.



Backup — Z7 · slow_fall_events — long- τ _fall study (PDS2) (3/8)

figure: zip7_PDS2_slow_fall.png — event rows 4–6 of 10, left half of the channels


10 random events with $\tau_{\text{fall}} > 1.5$ ms drawn in every channel: 11 channels show normal fast pulses; the long fall is the PDS2 low-frequency hardware noise, not slow physics.



Backup — Z7 · slow_fall_events — long- τ _fall study (PDS2) (4/8)

figure: zip7_PDS2_slow_fall.png — event rows 4–6 of 10, right half of the channels


10 random events with $\tau_{\text{fall}} > 1.5$ ms drawn in every channel: 11 channels show normal fast pulses; the long fall is the PDS2 low-frequency hardware noise, not slow physics.



Backup — Z7 · slow_fall_events — long- τ _fall study (PDS2) (5/8)

figure: zip7_PDS2_slow_fall.png — event rows 7–9 of 10, left half of the channels


10 random events with $\tau_{\text{fall}} > 1.5$ ms drawn in every channel: 11 channels show normal fast pulses; the long fall is the PDS2 low-frequency hardware noise, not slow physics.



Backup — Z7 · slow_fall_events — long- τ _fall study (PDS2) (6/8)

figure: zip7_PDS2_slow_fall.png — event rows 7–9 of 10, right half of the channels


10 random events with $\tau_{\text{fall}} > 1.5$ ms drawn in every channel: 11 channels show normal fast pulses; the long fall is the PDS2 low-frequency hardware noise, not slow physics.



Backup — Z7 · slow_fall_events — long- τ _fall study (PDS2) (7/8)

figure: zip7_PDS2_slow_fall.png — event rows 10–10 of 10, left half of the channels


10 random events with $\tau_{\text{fall}} > 1.5$ ms drawn in every channel: 11 channels show normal fast pulses; the long fall is the PDS2 low-frequency hardware noise, not slow physics.



Backup — Z7 · slow_fall_events — long- τ _fall study (PDS2) (8/8)

figure: zip7_PDS2_slow_fall.png — event rows 10–10 of 10, right half of the channels


10 random events with $\tau_{\text{fall}} > 1.5$ ms drawn in every channel: 11 channels show normal fast pulses; the long fall is the PDS2 low-frequency hardware noise, not slow physics.



Backup — Z7 · deliverables/nxm — PCA templates nxm0–4 (1/6)

figure: zip7_pca_templates.png — channel panels 1–2 of 11, top to bottom


nxm0 = mean curve; nxm1–4 = PCA components (oscillating basis vectors, may be negative), all peak-normalized; PC1+PC2 capture 96–98% of the shape variance. PDS2 replaced by the PDS1 template in the delivered files.



Backup — Z7 · deliverables/nxm — PCA templates nxm0–4 (2/6)

figure: zip7_pca_templates.png — channel panels 3–4 of 11, top to bottom


nxm0 = mean curve; nxm1–4 = PCA components (oscillating basis vectors, may be negative), all peak-normalized; PC1+PC2 capture 96–98% of the shape variance. PDS2 replaced by the PDS1 template in the delivered files.



Backup — Z7 · deliverables/nxm — PCA templates nxm0–4 (3/6)

figure: zip7_pca_templates.png — channel panels 5–6 of 11, top to bottom


nxm0 = mean curve; nxm1–4 = PCA components (oscillating basis vectors, may be negative), all peak-normalized; PC1+PC2 capture 96–98% of the shape variance. PDS2 replaced by the PDS1 template in the delivered files.



Backup — Z7 · deliverables/nxm — PCA templates nxm0–4 (4/6)

figure: zip7_pca_templates.png — channel panels 7–8 of 11, top to bottom


nxm0 = mean curve; nxm1–4 = PCA components (oscillating basis vectors, may be negative), all peak-normalized; PC1+PC2 capture 96–98% of the shape variance. PDS2 replaced by the PDS1 template in the delivered files.



Backup — Z7 · deliverables/nxm — PCA templates nxm0–4 (5/6)

figure: zip7_pca_templates.png — channel panels 9–10 of 11, top to bottom


nxm0 = mean curve; nxm1–4 = PCA components (oscillating basis vectors, may be negative), all peak-normalized; PC1+PC2 capture 96–98% of the shape variance. PDS2 replaced by the PDS1 template in the delivered files.



Backup — Z7 · deliverables/nxm — PCA templates nxm0–4 (6/6)

figure: zip7_pca_templates.png — channel panels 11–11 of 11, top to bottom

nxm0 = mean curve; nxm1–4 = PCA components (oscillating basis vectors, may be negative), all peak-normalized; PC1+PC2 capture 96–98% of the shape variance. PDS2 replaced by the PDS1 template in the delivered files.




Backup — Z7 · 1x1 summed templates PT / PS1 / PS2

figures: deliverables/1x1/plots/{PT,PS1,PS2}/zip7_*.png


Single-template (1x1) summed curves: peak-normalized average of the per-channel nxm0 templates — PT = all channels, PS1 / PS2 = side-1 / side-2 only.

processing: raw MIDAS trace -> 100 kHz 4th-order Butterworth low-pass (scipy sosfilt, steady-state init)
-> baseline subtract (median of samples 2000-12000) -> normalize by GLOBAL trace peak (no peak window)
-> 2-exp fit $y = A(\exp(-t/t_{\text{fall}}) - \exp(-t/t_{\text{rise}}))$, ALL 5 params free incl. pretrigger (curve_fit, pretrigger in 16050-3000, maxfev=10000)
-> fit_ok := amp>0 and 0<t_rise<fall -> NRMSE := RMS(fit residual)/fitted pulse peak, RECORDED ONLY (no cut)
-> align := shift the MEASURED LP trace by (fitted pretrigger - 16050), sub-sample up.interp
this figure: PT 1x1 template of zip7: peak-normalized average of the per-channel nxm0 PCA templates (all channels); channel nxm0 = mean of the clean fitted-curve population (fit_ok, NRMSE<=0.4, t_rise<=0.3ms)




Z7 PT

processing: raw MIDAS trace -> 100 kHz 4th-order Butterworth low-pass (scipy sosfilt, steady-state init)
-> baseline subtract (median of samples 2000-12000) -> normalize by GLOBAL trace peak (no peak window)
-> 2-exp fit $y = A(\exp(-t/t_{\text{fall}}) - \exp(-t/t_{\text{rise}}))$, ALL 5 params free incl. pretrigger (curve_fit, pretrigger in 16050-3000, maxfev=10000)
-> fit_ok := amp>0 and 0<t_rise<fall -> NRMSE := RMS(fit residual)/fitted pulse peak, RECORDED ONLY (no cut)
-> align := shift the MEASURED LP trace by (fitted pretrigger - 16050), sub-sample up.interp
this figure: PS1 1x1 template of zip7: peak-normalized average of the per-channel nxm0 PCA templates (side-1 channels); channel nxm0 = mean of the clean fitted-curve population (fit_ok, NRMSE<=0.4, t_rise<=0.3ms)



Z7 PS1

processing: raw MIDAS trace -> 100 kHz 4th-order Butterworth low-pass (scipy sosfilt, steady-state init)
-> baseline subtract (median of samples 2000-12000) -> normalize by GLOBAL trace peak (no peak window)
-> 2-exp fit $y = A(\exp(-t/t_{\text{fall}}) - \exp(-t/t_{\text{rise}}))$, ALL 5 params free incl. pretrigger (curve_fit, pretrigger in 16050-3000, maxfev=10000)
-> fit_ok := amp>0 and 0<t_rise<fall -> NRMSE := RMS(fit residual)/fitted pulse peak, RECORDED ONLY (no cut)
-> align := shift the MEASURED LP trace by (fitted pretrigger - 16050), sub-sample up.interp
this figure: PS2 1x1 template of zip7: peak-normalized average of the per-channel nxm0 PCA templates (side-2 channels); channel nxm0 = mean of the clean fitted-curve population (fit_ok, NRMSE<=0.4, t_rise<=0.3ms)



Z7 PS2