

Speaker Script / 演讲稿 — SNOLAB R4 Phonon Pulse Templates

约 9 分钟，11 页 + 1 页 backup（弱探测器 Z22 示例，不占正片时间）。每页先中文、后英文，内容一一对应，**英文可以直接照读**。

原则：稿子是说的话，不是幻灯片的复读——具体参数都在屏幕上，嘴里讲思路和为什么。

斜体是给自己的提示，不用念。

Slide 1 — Title (约 30 秒)

中文：大家下午好。今天讲我为 SNOLAB Run 4 做的声子脉冲模板。简单说，就是给 13 个探测器的每个声子通道做出一个“标准脉冲形状”，给 optimal filter 用。我按实际做的顺序讲：事件怎么选的，波形怎么拟合和对齐的，两个质量 cut 是怎么从数据里定出来的，最后是交付的两族模板。

English: Good afternoon everyone. Today I'll talk about the phonon pulse templates I built for SNOLAB Run 4. In short: for each phonon channel of thirteen detectors, I build a standard pulse shape for the optimal filter to use. I'll go in the order I actually did the work: how the events were selected, how each trace is fitted and aligned, how the two quality cuts were derived from the data itself, and finally the two template families I delivered.

Slide 2 — From the Ge-activation K-line to an event sample (约 65 秒)

中文：先说事件从哪来。做模板需要一大批干净、彼此相同的脉冲。我用的是 Ge 活化 K 线：Cf 活化之后，每个探测器里都会持续出现 10.37 keV 的单能事件，这是现成的理想样本。屏幕上 13 个探测器的 PTOFamps 谱——可以把 PTOFamps 理解成官方处理链给出的能量代理量。每格里的红线，是 Saab 教授的分析标出来的 K 线位置，我直接拿来用。我做的事情非常简单：以红线为中心开一个窗，上下各 1.35 倍——这个倍数是粗略目测定的，带点随意性，反正开得够宽就行——窗内的事件全要，不加任何别的条件。然后把每个选中事件所有通道的原始波形，原封不动存下来，一共约 120 GB。大家看这些谱：安静的探测器比如 Z7，K 线峰和噪声峰分得很开；但不少噪声大的探测器，红线就插在噪声堆里——窗口一开，噪声肯定跟着进来。这是故意的：我宁愿先多收，也要让后面每一步筛选清清楚楚、随时可以回退——因为存的是原始数据，什么都没丢。

English: Let me start with where the events come from. To build a template, you need a large sample of clean, identical pulses. I use the germanium activation K-line: after the californium activation, every detector keeps producing mono-energetic events at 10.37 keV — a perfect, ready-made sample. What you see here are the PTOFamps spectra of all thirteen detectors — you can think of PTOFamps as the energy estimate from the official processing. The red line in each panel is the K-line position marked by Professor Saab's analysis, and I simply take that as given. What I do is very simple: open a window around the red line, a factor of one point three five on each side — a rough, somewhat arbitrary, eyeballed choice; it just needs to be wide enough — keep everything inside, and apply no other condition. Then, for every selected event, I save the raw traces of all

channels, completely untouched — about a hundred and twenty gigabytes in total. Now look at the spectra: on a quiet detector like Z7, the K-line peak is well separated from the noise peak. But on the weaker detectors, the red line sits right inside the noise population — so the window will let noise in. And that's deliberate. I'd rather collect too much, and keep every later selection explicit and reversible — nothing is lost, because what I stored is raw.

Slide 3 — Fit quality at a glance / data overview (约 40 秒)

中文：先看单个事件层面，拟合到底靠不靠谱。网格图里一行是一个事件、一列是一个通道，蓝色是滤波后的波形，红色是拟合。规律非常清楚：左边这批是噪声触发，它在**所有**通道里同时拟合失败；右边是真正的 K 线事件，**所有**通道同时拟合得很好，残差只有百分之几。也就是说，事件干干净净地分成两类——要么全通道都好，要么全通道都坏。至于每一条波形具体怎么滤波、怎么拟合、怎么对齐，就是下一页的内容。

English: First, a sanity check at the level of individual events. In these grids, each row is one event and each column is one channel — blue is the filtered trace, red is the fit. The pattern is very clear. The left block is noise triggers: the fit fails in **every** channel at the same time. The right block is genuine K-line events: **every** channel fits well at the same time, with a residual of only a few percent. So the events split cleanly into two kinds — good in every channel, or bad in every channel. Exactly how each trace is filtered, fitted and aligned is the next slide.

Slide 4 — Per-trace algorithm (约 80 秒)

中文：拿到原始波形后，每一条都过同样的五步。完整参数都在屏幕上，我挑重点讲。前三步是准备：低通滤波抹掉高频噪声，扣基线，按峰值归一。核心是第四步的拟合：一个双指数函数，一个指数管上升、一个管下降——这是声子脉冲的标准形状。重点在于，五个拟合参数里，连 pretrigger——脉冲的起始时刻——也是放开的。为什么？因为真实触发时刻每个事件都不一样——拟合结果显示，pretrigger 普遍比名义位置晚两百多个采样点。如果把它钉死，这个误差就会被转嫁到上升、下降时间上，把所有参数都带歪。拟合完，每条波形记两个质量指标：一个叫 fit_ok，检查参数物不物理——幅度得是正的，上升得比下降快；另一个叫 NRMSE，全称 Normalized Root-Mean-Square Error，归一化均方根误差——就是拟合残差除以脉冲峰值，衡量拟合得好不好；NRMSE 等于 0.05，意思就是拟合偏差只有脉冲高度的 5%。注意，这一步只记录、不筛选。第五步——对齐，以及对齐后的效果，下一页单独讲。

English: Once I have the raw traces, every single one goes through the same five steps — the full parameters are on the slide, so let me just walk through the ideas. The first three steps are preparation: a low-pass filter to smooth away high-frequency noise, baseline subtraction, and normalization to the peak. The core is the fit: a two-exponential function — one exponential for the rise, one for the fall — which is the standard shape of a phonon pulse. And here's the important part: among the five fit parameters, even the pretrigger — the pulse start time — is left free. Why? Because the true trigger time is different for every event — the fits tell us the pretrigger is typically more than two hundred samples later than the nominal position. If you pin it, that error gets

pushed into the rise and fall times, and everything comes out biased. After the fit, each trace gets two quality numbers. One is called `fit_ok` — a sanity check that the parameters are physical: positive amplitude, rise faster than fall. The other is the NRMSE — the normalized root-mean-square error — which is simply the fit residual divided by the pulse height, so it tells you how good the fit actually is: an NRMSE of zero point zero five means the fit misses the data by five percent of the pulse height. At this stage I only record these; I don't cut on them yet. Step five — the alignment, and what it produces — gets its own slide next.

Slide 5 — Alignment result (约 40 秒)

中文：这一页讲第五步——对齐。每一条实测波形，按它自己拟合出的 `pretrigger` 减去名义的 16050 平移——纯平移，波形本身没有任何再生成。图里那条竖虚线，就是所有波形对齐到的公共 `pretrigger`。对齐之后，把所有拟合曲线画在这个公共起点上，就是这两张扇形图。左边是全部物理拟合的曲线——可以看到主束旁边有一些明显散开、下降很慢的曲线；右边是用 NRMSE 筛过之后，只剩一束非常紧、非常一致的形状。正是把触发抖动对齐掉、再筛掉坏拟合之后，剩下这条紧致的束，让我能做出一个良定义的模板。至于 NRMSE 这个筛选阈值具体怎么定的，下一页讲。

English: This slide is step five — alignment. Each measured trace is shifted by its own fitted `pretrigger` minus 16050 — a pure translation, nothing about the waveform is regenerated. The dotted vertical line is the common `pretrigger` everything is aligned to. Once aligned, I draw all the fitted curves from that common starting point — these two fan plots. On the left, all the physical fits: you can see some clearly separated, slowly-falling stragglers spreading off the main bundle. On the right, after selecting on NRMSE, only one very tight, consistent shape family remains. Alignment removes the jitter, the NRMSE cut removes the bad fits, and what's left is this tight bundle — tight enough to build a template from. Exactly how the NRMSE threshold is set is the next slide.

Slide 6 — Quality cut 1: $\text{NRMSE} \leq 0.4$ (约 55 秒)

中文：现在讲第一个 cut 怎么定的。把所有物理拟合的 NRMSE 画成分布，每个探测器都是同一个样子：两个峰。左边的峰是拟合好的，典型值百分之五到百分之十；右边的峰在 1 到 2 附近，就是那些根本没有脉冲的噪声触发。两峰之间有一个很深的谷，位置在 0.4 左右——阈值就取 0.4。我想强调：这个数不是调出来的，是分布自己长出来的。图里是 Z7 的 PBS1——全场用的就是这条通道——两个峰离得很远，谷一目了然。较弱的探测器是同样的双峰图景，只是噪声峰是主体，对它们来说这一刀就是把真脉冲从噪声堆里捞出来的那一刀；具体例子放在 backup 页。

English: Now, how the first cut was set. If you histogram the NRMSE of all physical fits, every detector shows the same picture: two populations. The left one is the good fits, typically five to ten percent. The right one, around one or two, is the noise triggers — traces with no pulse in them at all. In between there's a deep valley, at about zero point four — and that's where I put the threshold. I want to stress: this number wasn't tuned, it's simply what the distribution itself shows. The plot is Z7 PBS1 — the same channel used throughout the talk — where the two populations are far apart and the valley is obvious. The weak detectors show the same bimodal picture with the

noise peak dominating, and for them this cut is exactly what pulls the real pulses out of the noise; an example is in the backup slide.

Slide 7 — The rejected population is noise (约 45 秒)

中文：定了 cut 还不够，得验证它切掉的确实是垃圾。左边是被切掉的事件的**原始**波形——可以看到，里面就没有脉冲，纯粹是噪声，只不过一条很慢的曲线恰好凑在了它上面。右边换个角度看同一件事，这张图有三层：底下灰蓝色的是对齐后的实测波形，上面绿色是**通过** cut 的拟合曲线，红色是**被切掉的**；右上角标着两个群体各自的中位 NRMSE 和事件数。可以看到绿的又快又一致，聚成一束；红的四处发散，中位 NRMSE 是绿的四十倍——1.87 对 0.045。在 Z7 这条通道上，被切掉的只有 77 个，不到百分之四。结论很干净：这一刀切掉的是噪声，没有伤到物理。

English: Setting a cut isn't enough — I have to verify that what it removes really is junk. On the left are the **raw** traces of the rejected events. As you can see, there is simply no pulse there — it's pure noise, on which some slow curve happened to converge. On the right, the same thing from another angle — this figure has three layers: the gray-blue underneath is the aligned measured traces, green on top is the fitted curves that **pass** the cut, and red is the ones **removed**; the top-right corner shows each population's median NRMSE and its event count. The green ones are fast and consistent, bundled together; the red ones scatter everywhere, with a median NRMSE about forty times higher — one point eight seven against zero point zero four five. On this Z7 channel, only seventy-seven events are removed, under four percent. So the conclusion is clean: this cut removes noise, and it doesn't touch the physics.

Slide 8 — Follow-up 1: the slow-fall tail (约 50 秒)

中文：第一条线索来自 cut 之后的扇形图：还剩一些“拖尾巴”的曲线，下降特别慢。我的办法是抽样验证：在通过 0.4 的事件里，挑下降时间超过 1.5 毫秒的，随机抽 10 个，把每个事件在全部 12 个通道里的原始波形和拟合画出来对比。结论有两层。第一层：**抽到的这些事件都是真脉冲**——所以保留，我没有为下降时间设任何 cut。第二层，看左边这张图——我特别想强调这一点：**Z7 本身是最好的探测器，但它有一个坏通道，PDS2**。你看每一行是一个事件，除了 PDS2 那一列，其他每一列都是干干净净的快脉冲；唯独 PDS2 这一列，波形上叠着一个大幅度的低频晃动，拟合去追那个晃动，才给出了几毫秒的假“下降时间”。也就是说，慢的不是事件，是 PDS2 这一个通道坏。右边这两张图是同一件事的量化，横轴都是 0 到 20 毫秒，可以直接比：上面是正常通道 PAS1——所有事件挤在 0.25 毫秒一根窄峰里，1 毫秒之后什么都没有；下面是 PDS2——一条明显的宽尾一路拖到五六毫秒，中位数 0.51 毫秒，是正常通道的两倍。所以最夸张的拖尾不是慢物理，是单通道的低频伪影，而且这个探测器好不好、跟这一个通道坏不坏是两回事。

English: The first lead comes from the fan plot after the cut: some curves still have long tails — a very slow fall. My approach was to verify by sampling: among the events that pass the cut, take the ones with a fall time above one and a half milliseconds, randomly sample ten of them, and draw raw versus fit in every one of the twelve channels. The conclusion has two layers. First: **the sampled events are real pulses** — so they stay in, and I apply no fall-time cut at all. Second —

and this is the point I really want to make with the left plot: **Z7 is the best detector overall, but it has one bad channel, PDS2**. Each row is one event, and in every column except PDS2 you see a clean, fast pulse; only in the PDS2 column is there a large low-frequency swing riding on top of the trace, and the fit chases that swing, which is what produces the fake fall times of several milliseconds. So it's not the event that's slow — it's that one channel, PDS2, that's misbehaving. The two plots on the right are the same story quantified, on the same zero-to-twenty-millisecond axis so you can compare directly: the top one is a normal channel, PAS1 — everything piled into a narrow spike at 0.25 milliseconds, nothing beyond one; the bottom one is PDS2 — a clear broad tail stretching out to five or six milliseconds, with a median of 0.51 milliseconds, twice that of the normal channel. So the most extreme tails are not slow physics; they're a one-channel low-frequency artifact — and a detector being good overall and one of its channels being bad are two separate things.

Slide 9 — Template family 1: 2-exp weighted (1x1) (约 50 秒)

中文：现在到产出。第一族模板走解析路线：每个通道，把所有物理拟合的曲线放到同一个 pretrigger 位置，做加权平均——权重是 NRMSE 平方的倒数。这个设计的好处是：拟合差的事件权重自动变得极小，等于被压没了，但我不需要人为剔除任何事件。屏幕上就是参与平均的这束曲线。因为平均的对象是解析函数，做出来的模板天然光滑、完全没有噪声。这就是标准的单模板，1x1。每个通道一条，交付成峰值归一的 32768 点 ROOT 直方图，另外还有把各通道求和得到的 PT、PS1、PS2 模板。

English: Now to the deliverables. The first template family takes the analytic route: for each channel, I put all the physical fitted curves at the same pretrigger and take a weighted average — the weight is one over NRMSE squared. The nice property of this design is that badly-fit events automatically get a vanishingly small weight, so they're effectively suppressed — but I never have to remove anything by hand. What's on screen is the bundle of curves that goes into that average. And because I'm averaging analytic functions, the resulting template is smooth and completely noise-free by construction. That's the standard single template — the one-by-one — delivered as a peak-normalized 32768-bin ROOT histogram per channel, plus the summed PT, PS1 and PS2 templates.

Slide 10 — Template family 2: NxM PCA (约 55 秒)

中文：第二族模板是为了捕捉不同事件之间脉冲形状的变化。做法是对干净的拟合曲线做主成分分析，也就是 PCA——输入是通过 NRMSE cut、再加一条上升时间上限（挡掉缓慢的基线漂移）之后剩下的曲线。直观理解：黑色的 nxm0 是平均形状；后面四条彩色的，是数据里最主要的四个“变形方向”——比如上升更慢一点、下降更快一点。真实脉冲就用这五条的线性组合去拟合。效果非常好：前两个成分就已经覆盖了 96 到 98% 的形状差异。也就是说，事件之间的形状差异被显式地表示进了模板空间里。最后一步，交付之前把五条模板统一归一到峰值为 1，方便对比和使用——这就是最终产物。

English: The second family is built to capture the variation in pulse shape across events. I run a principal component analysis — PCA — on the clean fitted curves: the ones passing the NRMSE cut plus a rise-time ceiling that removes slow baseline drift. The intuition: the black curve, nxm-zero, is the average shape; the four colored ones are the four main "directions of deformation" in the data — say, a slightly slower rise, or a faster fall. A real pulse is then fitted as a linear combination of these five. And it works remarkably well: the first two components already cover ninety-six to ninety-eight percent of the shape variation. So the shape variation across events is explicitly represented inside the template space. As a final step before delivery, all five templates are normalized to unit peak, so they're easy to compare and use — and that is the final product.

Slide 11 — Delivered, and what remains (约 30 秒)

中文：一句话总结：我为全部 13 个探测器做出了两族模板——1x1 和 NxM PCA，都按官方格式交付了。两个 cut 都是从数据里读出来、又回到原始波形验证过的。谢谢大家，欢迎提问。（如被问后续：下一步是对新模板跑一遍组里的验证流程。）

English: In one line: I built two template families for all thirteen detectors — the 1x1 and the NxM PCA — and delivered both in the official format. Both cuts were read off the data and verified in the raw traces. Thank you — happy to take questions. (If asked about next steps: run the group's validation procedure on the new templates.)

Backup — Weak detectors (仅备查，不占正片时间)

中文：弱探测器（Z1/Z4/Z6/Z18/Z19/Z22/Z24）的 K 线插在噪声堆里，窗口收进来的是噪声为主的混合体。NRMSE 分布同样双峰、但噪声峰是主体，同一条 0.4 的 cut 把真脉冲捞出来——Z22 PCS1：保留 7198、切掉 4788，约四成。保留束比安静探测器宽；红色里那些陡峭曲线是拟合抓到噪声尖刺，不是快脉冲被切掉。 $\tau_{\text{rise}} \leq 0.3 \text{ ms}$ 在弱探测器上再切掉 NRMSE 之后余下的 55–74%（Z22 是 71%），主要是残余的慢基线漂移。

English: On the weak detectors (Z1/Z4/Z6/Z18/Z19/Z22/Z24) the K-line sits inside the noise population, so the window admits a noise-dominated mixture. The NRMSE distribution is still bimodal, only with the noise peak dominating, and the same 0.4 cut digs the real pulses out — on Z22 PCS1, 7198 kept and 4788 cut, about forty percent. The kept bundle is broader than on a quiet detector, and the steep red curves are fits latching onto sharp noise spikes, not fast pulses being cut. The rise-time ceiling then removes another 55–74% of what survived the NRMSE cut on the weak zips (71% on Z22) — residual slow baseline drift.

备用问答提示 / Q&A hints

- $\tau_{\text{rise}} \leq 0.3 \text{ ms}$ 相对上一步切掉多少？ / How much does the rise-time ceiling remove, relative to the NRMSE step? 安静探测器几乎不切：Z7 1.6%（大头是坏通道 PDS2，单通道 21%，其余通道 $\leq 0.4\%$ ）、Z9 4.7%、Z16 5.9%。弱探测器切 55–74%：Z24 74%、Z4 和 Z22 各 71%、Z18 71%、Z19 68%、Z6 56%。13 个 zip 合计 54%。切掉的主要是通过 NRMSE 的残余慢漂移。 / Quiet detectors

lose almost nothing: Z7 1.6% (mostly the bad channel PDS2 at 21%; every other channel $\leq 0.4\%$), Z9 4.7%, Z16 5.9%. Weak detectors lose 55–74% (Z24 74%, Z4/Z22 71%, Z18 71%, Z19 68%, Z6 56%). Pooled over all 13 zips: 54%. What it removes is residual slow baseline drift that survived the NRMSE cut.

- **为什么不用实测波形的平均做模板？ / Why not average the measured traces?** 解析拟合曲线本身没有噪声，加权平均出来的模板天然光滑；NRMSE 权重让拟合差的事件自动降权。对齐叠加图上实测均值（红）和加权拟合均值（橙）可以互相对照。 / The analytic fitted curves carry no noise, so their weighted mean is smooth by construction; the NRMSE weight down-weights badly-fit events automatically. The overlays show the measured mean (red) and the weighted fitted mean (orange) side by side for comparison.

- **0.4 怎么定的，敏感吗？ / How was 0.4 chosen — is it sensitive?** 取在双峰分布的谷底，谷里本来就没什么事件——这正是选谷底的理由。还有交互 notebook，改阈值几秒内能重跑所有探测器。 / It sits in the valley of the bimodal distribution, where there are few events by construction — that's exactly why the valley was chosen. An interactive notebook re-runs any threshold within seconds.

- **两族模板的 cut 为什么不一样？ / Why do the two families apply cuts differently?** 1x1 不硬切，靠权重：好拟合 NRMSE ≈ 0.05 权重 400，噪声 NRMSE ≈ 2 权重 0.25，差四个数量级，噪声自动压没，不需要人为定一刀。PCA 对离群点敏感——一条噪声曲线就能污染主成分方向——所以 NxM 的输入必须硬切干净（ $\text{fit_ok} + \text{NRMSE} \leq 0.4 + \tau_{\text{rise}} \leq 0.3 \text{ ms}$ ）。 / The 1x1 uses weighting instead of a hard cut: a good fit at NRMSE 0.05 gets weight 400, a noise trigger at 2 gets 0.25 — four orders of magnitude apart, so noise is automatically suppressed without choosing a threshold. PCA is sensitive to outliers — one noise curve can contaminate the components — so the NxM input must be hard-cut clean.

- **慢脉冲群体到底留不留？ / Keep the slow pulses or not?** 这是明确记录在案的待定事项。目前 PCA 的输入应用了 τ_{rise} cut；这个取舍要不要放开，正是想和大家讨论决定的。 / This is the documented open item. The PCA input currently applies the rise-time cut; whether to relax that trade-off is exactly what I'd like to discuss and decide.

术语自查表（只给自己，不进 PPT）

过一遍下面每一行；哪一条还觉得说不清楚，提前告诉我，我把它从幻灯片里去掉或换个说法。

术语	一句话解释
zip / Z7	一个探测器（一块 Ge 晶体）；Run 4 一共 13 个，Z7 是最安静的一个
声子通道 PAS1...PES2	每个探测器的声子读出通道，S1/S2 是晶体的两个面；Z7 的 PFS2 坏了所以只有 11 个
series	一段连续采数（约 1–2 小时），文件按 series 组织；每个探测器 27–30 个
MIDAS raw	DAQ 写出的原始波形数据格式，完全未经处理
PTOFamps	官方处理链里"总声子 optimal filter 幅度"，可当作能量的代理量

术语	一句话解释
K 线	Cf 活化后 Ge 探测器出现的 10.37 keV K 壳层活化线，当单能参考源用
optimal filter (OF)	用"模板形状 + 噪声谱"对脉冲做加权拟合来估计幅度的标准方法
1x1 / NxM	1x1 = 每通道一个模板；NxM = 每通道多个模板（均值 + 主成分），能拟合形状变化
Butterworth 低通	一种平滑滤波器；100 kHz 截止把高频噪声滤掉；"4 阶"指截止后压制的陡峭程度
steady-state init	滤波器初始条件取稳态值，避免波形开头出现人为的启动瞬态
2-exp 模型	上升一个指数、下降一个指数： $y = A(e^{-(t/\tau_{\text{fall}})} - e^{-(t/\tau_{\text{rise}})})$ ；5 个自由参数
pretrigger / onset 16050	波形里名义的触发位置（第 16050 个采样点）；我把它当自由参数拟合
NRMSE	拟合残差的 RMS ÷ 拟合脉冲峰值；无量纲的"拟合好坏"指标
fit_ok	物理性检查：幅度为正且 $\tau_{\text{rise}} < \tau_{\text{fall}}$
亚采样插值对齐	平移量不是整数个采样点时用线性插值（np.interp）实现的平移
PCA	主成分分析：把一堆曲线分解成"均值 + 按方差大小排序的正交变化方向"
explained variance	每个主成分解释的形状方差占比；PC1+PC2 = 96–98%
TH1D	ROOT 软件的一维直方图对象；模板的交付格式（32768 个 bin）
cdmsbats / PulseTemplates	官方处理软件 / 它读取模板的配置目录和文件命名规范
PT / PS1 / PS2	求和模板：全部通道之和、S1 面之和、S2 面之和（各自峰值归一）